

Assignment 4

Solution (a)

$$C \approx \gamma 2^{A/2}$$

$$\Rightarrow \ln(C) \approx \ln(\gamma) + A/2 * \ln(2)$$

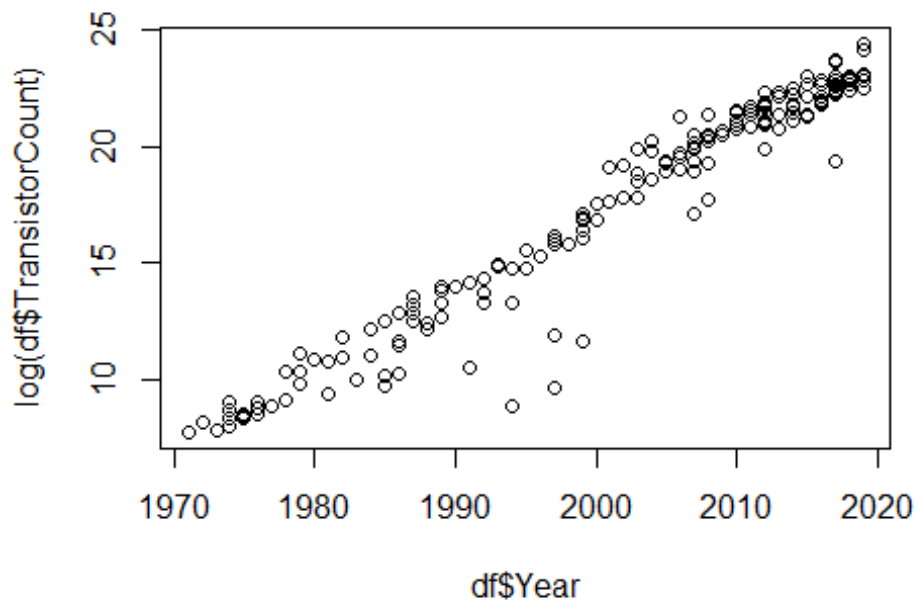
This equation is of the form $y \approx \beta_1 + \beta_2 x$

where $y = \ln(C)$, $x = A$, $\beta_1 = \ln(\gamma)$, $\beta_2 = \ln(2)/2$

The coefficient of A - β_2 is $\ln(2)/2 = 0.3465736$

Solution (b)

```
df = read.csv("C:\\temp\\mooreslawdata.csv",header = TRUE)
plot(df$Year, log(df$TransistorCount))
```



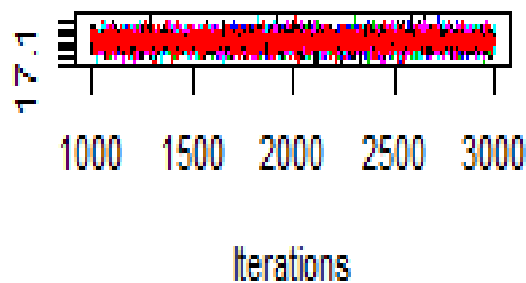
```
#add new columns for year and log of transistor count to be used in JAGS model
df$A = df$Year
df$C_log = log(df$TransistorCount)
```

Solution (c)(i)

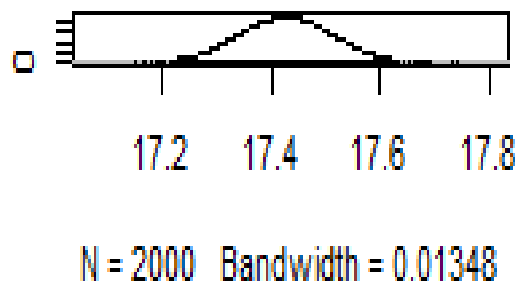
```
model {  
  for (j in 1:length(C_log)) {  
    C_log[j] ~ dnorm(beta1 + beta2*(A[j] - mean_A), sigmasqinv)  
  }  
  
  beta1 ~ dnorm(0,1/1000^2)  
  beta2 ~ dnorm(0,1/1000^2)  
  sigmasqinv ~ dgamma(0.001, 0.001)  
  sigmasq <- 1/sigmasqinv  
  mean_A <- mean(A)  
}  
  
inits <- list(list(beta1=10000, beta2=10000, sigmasqinv=0.01, .RNG.name="base::  
Wichmann-Hill", .RNG.seed=2),  
              list(beta1=10000, beta2=10000, sigmasqinv=0.00001, .RNG.name="base::  
Wichmann-Hill", .RNG.seed=3),  
              list(beta1=10000, beta2=-10000, sigmasqinv=0.01, .RNG.name="base  
::Wichmann-Hill", .RNG.seed=4),  
              list(beta1=10000, beta2=-10000, sigmasqinv=0.00001, .RNG.name="base  
::Wichmann-Hill", .RNG.seed=5),  
              list(beta1=-10000, beta2=10000, sigmasqinv=0.01, .RNG.name="base  
::Wichmann-Hill", .RNG.seed=6),  
              list(beta1=-10000, beta2=10000, sigmasqinv=0.00001, .RNG.name="base  
::Wichmann-Hill", .RNG.seed=7),  
              list(beta1=-10000, beta2=-10000, sigmasqinv=0.01, .RNG.name="base  
::Wichmann-Hill", .RNG.seed=8),  
              list(beta1=-10000, beta2=-10000, sigmasqinv=0.00001, .RNG.name="base  
::Wichmann-Hill", .RNG.seed=9)  
            )  
library(rjags)  
  
## Warning: package 'rjags' was built under R version 3.6.3  
## Loading required package: coda  
## Warning: package 'coda' was built under R version 3.6.3  
## Linked to JAGS 4.3.0  
## Loaded modules: basemod,bugs  
  
m1 <- jags.model("C:\\temp\\mooreslaw.bug", df, inits, n.chains=8)
```

```
## Warning in jags.model("C:\\temp\\mooreslaw.bug", df, inits, n.chains = 8):  
## Unused variable "Processor" in data  
  
## Warning in jags.model("C:\\temp\\mooreslaw.bug", df, inits, n.chains = 8):  
## Unused variable "Year" in data  
  
## Warning in jags.model("C:\\temp\\mooreslaw.bug", df, inits, n.chains = 8):  
## Unused variable "TransistorCount" in data  
  
## Compiling model graph  
##   Resolving undeclared variables  
##   Allocating nodes  
## Graph information:  
##   Observed stochastic nodes: 178  
##   Unobserved stochastic nodes: 3  
##   Total graph size: 516  
##  
## Initializing model  
  
update(m1, 1000) # burn-in  
x1 <- coda.samples(m1, c("beta1", "beta2", "sigmasq"), n.iter=2000)  
  
plot(x1)
```

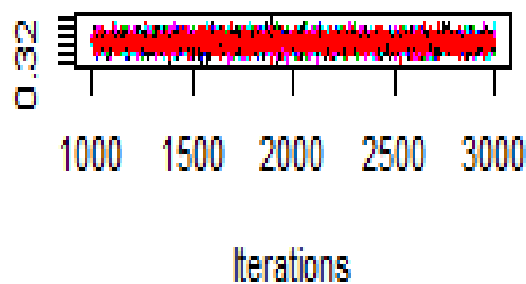
Trace of beta1



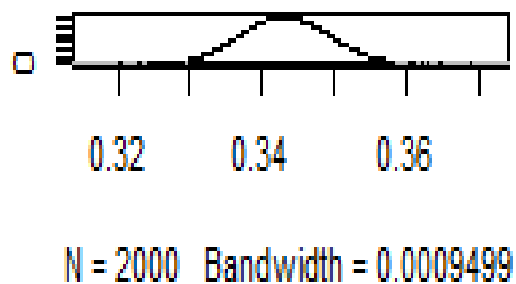
Density of beta1



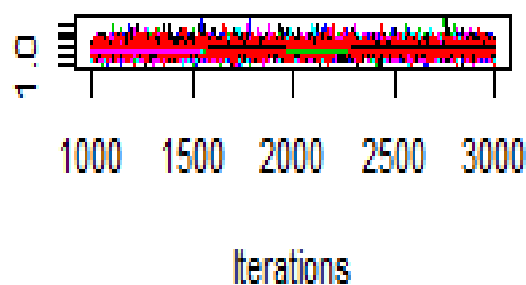
Trace of beta2



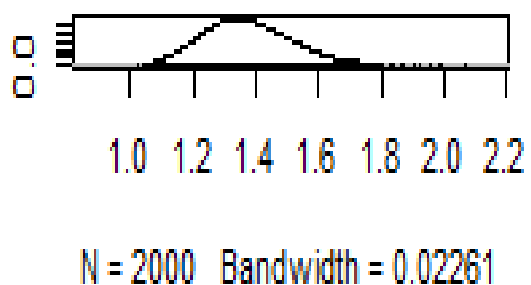
Density of beta2



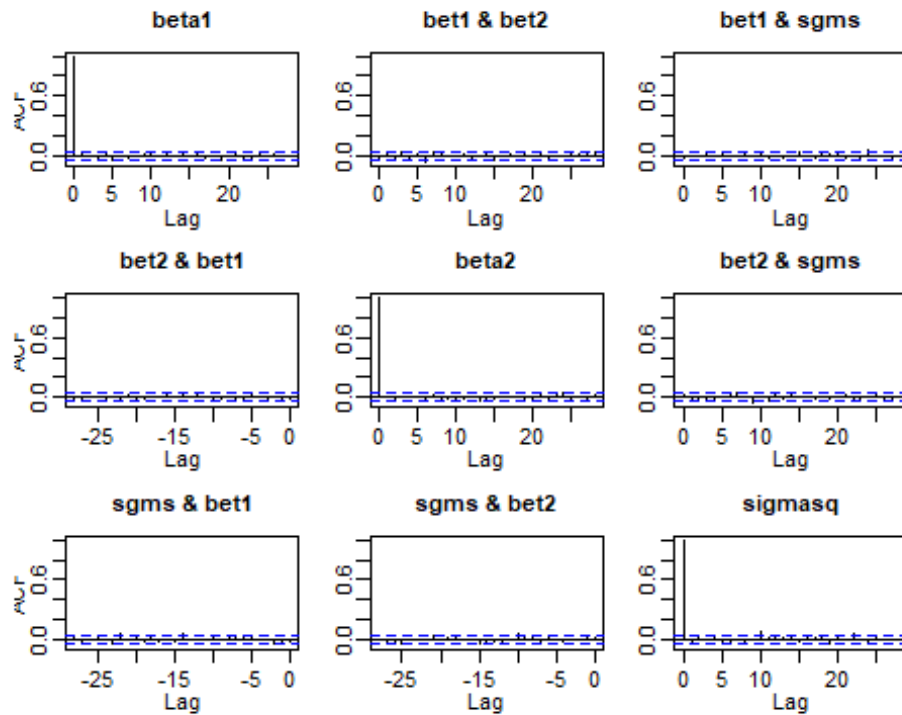
Trace of sigmasq



Density of sigmasq



```
acf(x1[[1]])
```



```
gelman.diag(x1, autoburnin = FALSE)
```

```
## Potential scale reduction factors:
##
##      Point est. Upper C.I.
## beta1          1          1
## beta2          1          1
## sigmasq        1          1
##
## Multivariate psrf
##
## 1
```

We can see from plots, chains are sampling from same regions, have indepdnent samples and gelman factors are 1 hence indicating convergence

Solution (c)(ii)

```
summary(x1)
```

```
##
## Iterations = 1001:3000
## Thinning interval = 1
## Number of chains = 8
## Sample size per chain = 2000
##
```

```
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##           Mean          SD Naive SE Time-series SE
## beta1    17.4258 0.088159 6.970e-04      6.910e-04
## beta2     0.3432 0.006212 4.911e-05      4.922e-05
## sigmasq   1.3774 0.148321 1.173e-03      1.185e-03
##
## 2. Quantiles for each variable:
##
##           2.5%    25%    50%    75%    97.5%
## beta1    17.2521 17.366 17.4259 17.4847 17.5982
## beta2     0.3311 0.339 0.3432 0.3474 0.3555
## sigmasq   1.1175 1.273 1.3666 1.4710 1.7020
```

Solution (c)(iii)

The slope is β_2 , as we can see from (c)(ii) posterior mean is 0.3432 The 95% posterior interval is (0.3311, 0.3555) and contains the value 0.3465736 from part a)

Solution (c)(iv)

The intercept is β_1 , and its posterior mean is 17.4258 The 95% posterior interval is (17.2521, 17.5982)

Solution (d)(i)

JAGS – MODEL

```
model {
  for (j in 1:length(C_log)) {
    C_log[j] ~ dnorm(beta1 + beta2*(A[j] - mean_A), sigmasqinv)
  }

  beta1 ~ dnorm(0,1/1000^2)
  beta2 ~ dnorm(0,1/1000^2)
  sigmasqinv ~ dgamma(0.001, 0.001)
  sigmasq <- 1/sigmasqinv
  mean_A <- mean(A)
  A_pred <- 2021
  C_pred_log ~ dnorm(beta1 + beta2*(A_pred - mean_A), sigmasqinv)
```

```
C_pred_linear <- exp(C_pred_log)
```

```
invention_year <- mean_A - beta1/beta2
```

```
}
```

```
m1 <- jags.model("C:\\temp\\mooreslawPred.bug", df, inits, n.chains=8, n.adapt=1000)
```

```
## Warning in jags.model("C:\\temp\\mooreslawPred.bug", df, inits, n.chains = 8, :
```

```
## Unused variable "Processor" in data
```

```
## Warning in jags.model("C:\\temp\\mooreslawPred.bug", df, inits, n.chains = 8, :
```

```
## Unused variable "Year" in data
```

```
## Warning in jags.model("C:\\temp\\mooreslawPred.bug", df, inits, n.chains = 8, :
```

```
## Unused variable "TransistorCount" in data
```

```
## Compiling model graph
```

```
##   Resolving undeclared variables
```

```
##   Allocating nodes
```

```
## Graph information:
```

```
##   Observed stochastic nodes: 178
```

```
##   Unobserved stochastic nodes: 4
```

```
##   Total graph size: 524
```

```
##
```

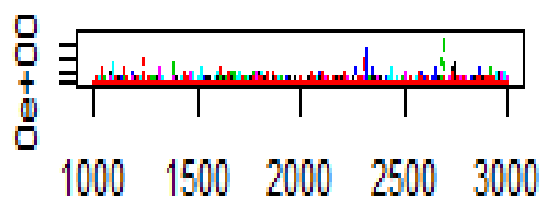
```
## Initializing model
```

```
update(m1, 1000) # burn-in
```

```
x2 <- coda.samples(m1, c("C_pred_log", "C_pred_linear", "invention_year"), n.iter=2000)
```

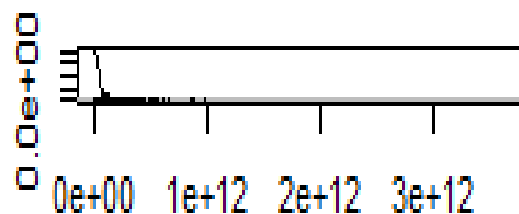
```
plot(x2)
```

Trace of C_pred_linear



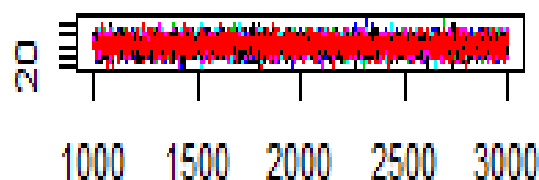
Iterations

Density of C_pred_linear



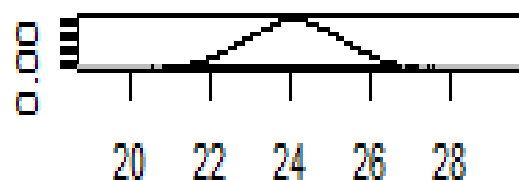
N = 2000 Bandwidth = 5.862e+09

Trace of C_pred_log



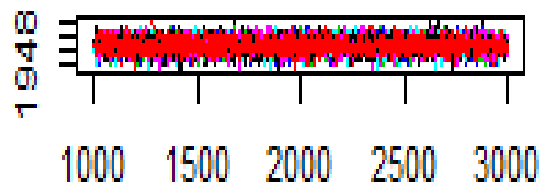
Iterations

Density of C_pred_log



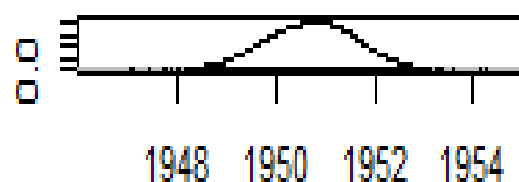
N = 2000 Bandwidth = 0.1792

Trace of invention_year



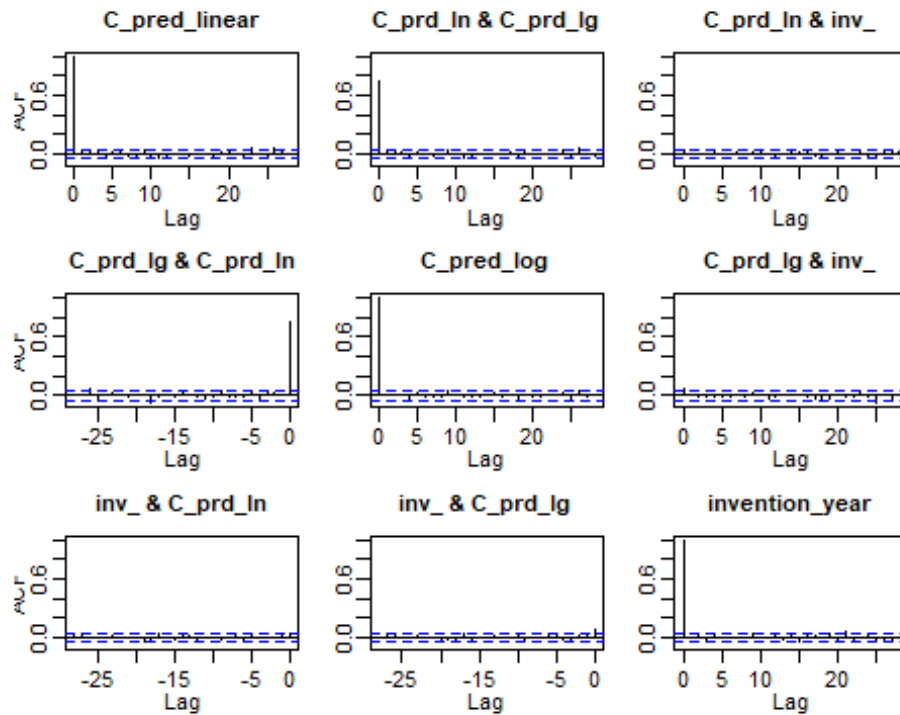
Iterations

Density of invention_year



N = 2000 Bandwidth = 0.1472


```
acf(x2[[1]])
```



```
gelman.diag(x2, autoburnin = FALSE)
```

```
## Potential scale reduction factors:
##
##           Point est. Upper C.I.
## C_pred_linear      1.01      1.01
## C_pred_log         1.00      1.00
## invention_year      1.00      1.00
##
## Multivariate psrf
##
## 1
```

We can see from plots, chains are sampling from same regions, have indepdnent samples and gelman factors are close to 1 hence indicating convergence

Solution (d)(ii)

```
summary(x2)
```

```
##
## Iterations = 1001:3000
## Thinning interval = 1
## Number of chains = 8
## Sample size per chain = 2000
##
```

```
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##              Mean          SD Naive SE Time-series SE
## C_pred_linear 5.802e+10 9.874e+10 7.806e+08      7.721e+08
## C_pred_log    2.410e+01 1.172e+00 9.263e-03      9.213e-03
## invention_year 1.951e+03 9.628e-01 7.612e-03      7.612e-03
##
## 2. Quantiles for each variable:
##
##              2.5%       25%       50%       75%       97.5%
## C_pred_linear 2.899e+09 1.345e+10 2.947e+10 6.482e+10 2.821e+11
## C_pred_log    2.179e+01 2.332e+01 2.411e+01 2.489e+01 2.637e+01
## invention_year 1.949e+03 1.950e+03 1.951e+03 1.951e+03 1.953e+03
```

Solution (d)(iii)

95% posterior predictive confidence interval for year 2021 is (2.899e+09,2.821e+11)

Solution (d)(iv)

The year indicated by $\bar{A} - \beta_1/\beta_2$ is the solution to the regression equation

$\log(C) \approx \beta_1 + \beta_2(A_i - \bar{A})$ when when $\log(C) = 0$ i.e. $C = 1$ (transistor count of 1)

The 95% posterior interval for invention year is (1949,1953)

Solution (e)(i)

```
m1 <- jags.model("C:\\temp\\mooreslaw.bug", df, inits, n.chains=8, n.adapt=1000)

## Warning in jags.model("C:\\temp\\mooreslaw.bug", df, inits, n.chains = 8,
## :
## Unused variable "Processor" in data

## Warning in jags.model("C:\\temp\\mooreslaw.bug", df, inits, n.chains = 8,
## :
## Unused variable "Year" in data

## Warning in jags.model("C:\\temp\\mooreslaw.bug", df, inits, n.chains = 8,
## :
## Unused variable "TransistorCount" in data

## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 178
##   Unobserved stochastic nodes: 3
##   Total graph size: 516
```

```
##
## Initializing model

df$adjusted_year = df$Year - mean(df$Year)
update(m1, 1000) # burn-in
x3 <- coda.samples(m1, c("beta1","beta2","sigma2"),n.iter=2000)

beta1_sim <- as.matrix(x3)[, "beta1"]
beta2_sim <- as.matrix(x3)[, "beta2"]

#combine beta1 and beta2 into a matrix
beta.sim <- cbind(beta1_sim,beta2_sim)

sigma2_sim <- as.matrix(x3)[, "sigma2"]
Nsim <- nrow(beta.sim)

mod <- lm(C_log ~adjusted_year, data=df)
X <- model.matrix(mod)

error.std.sim <- matrix(NA,Nsim,nrow(df))

#Simulated standard Error
for (s in 1:Nsim)
  error.std.sim[s,] <- (df$C_log - X %*% cbind(beta.sim[s,])) /sqrt(sigma2_sim[s])
```

Solution (e)(ii)

Replicate standard error is expected to have a standard normal distribution. Hence rnorm can be used to generate it.

```
#Simulate replicate Error from rnorm
error.rep.std.normal <-matrix(rnorm(Nsim*nrow(df)),Nsim,nrow(df))
```

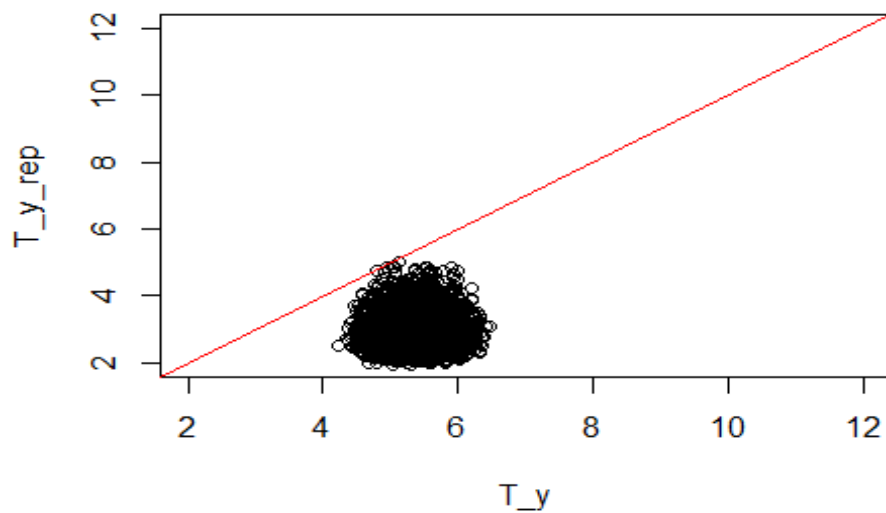
Solution (e)(iii)

Compute $T(y, X, \theta)$ and $T(y^{rep}, X, \theta)$

```
T_y <- apply(abs(error.std.sim), 1, max)
T_y_rep<- apply(abs(error.rep.std.normal), 1, max)
```

Solution (e)(iv)

```
plot(T_y,T_y_rep,xlim = c(2,12), ylim = c(2,12) )
abline(0,1,col ="red")
```



Solution (e)(v)

Compute p value for the outlier test statistic

```
mean(T_y_rep >= T_y)
```

```
## [1] 0.000125
```

Since p value is very low it indicates presence of outliers

Solution (e)(vi)

The extreme outlier can be defined as microprocessor that achieves the maximum difference for each simulation from standard normal error with highest probability

```
N_y <- apply(abs(abs(error.std.sim) - abs(error.rep.std.normal)), 1, which.max)
```

```
#Sort by highest incidence of a uP being maximum outlier wrt to the metric and print top 3
```

```
sort(table(N_y),decreasing=TRUE)[1:3]
```

```
## N_y
```

```
## 63 55 66
```

```
## 9730 5900 347
```

Index no 63 (F21 processor) has the highest occurrence (9730) out of 16000 simulations as extreme i.e. probability .608 for it bring the extreme outlier